




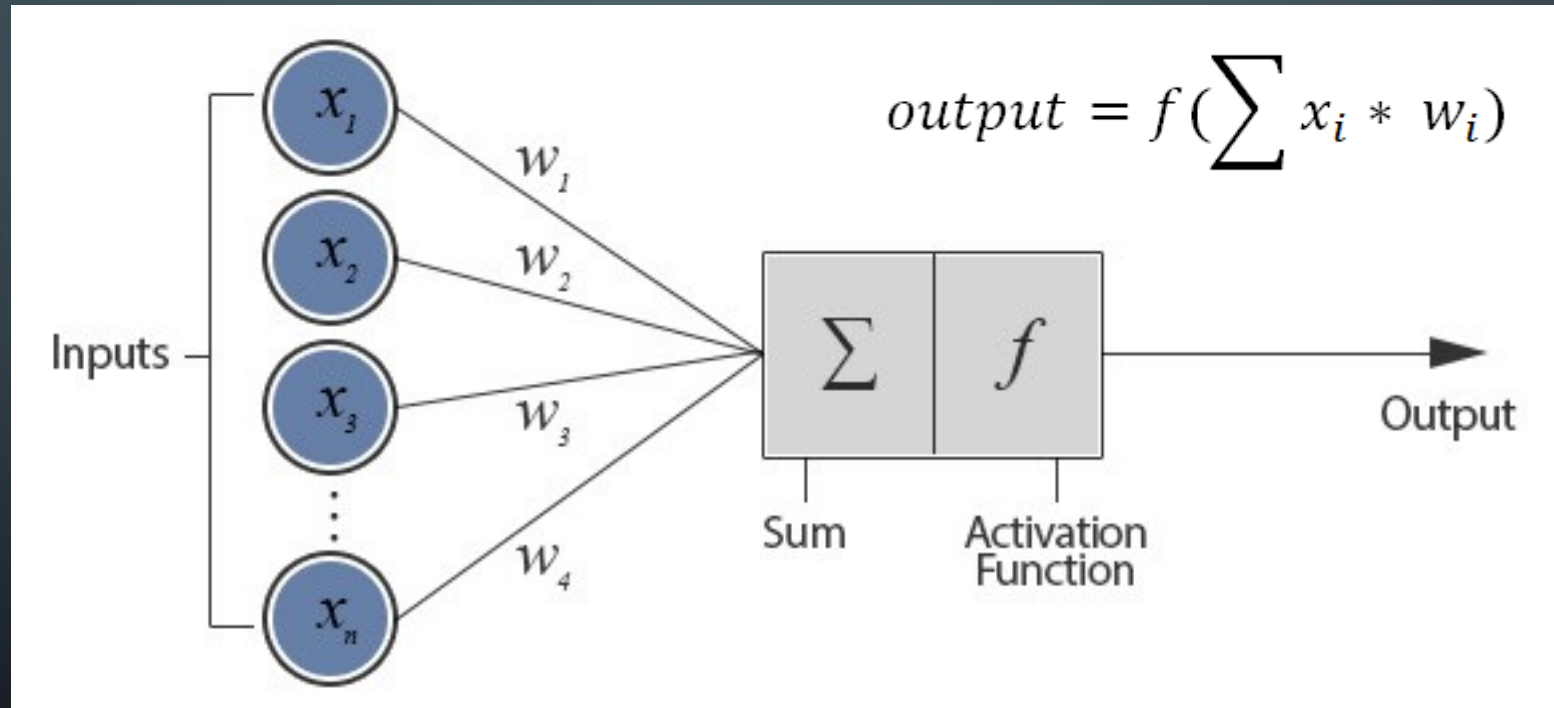
NEURAL NETWORKS

MOHAMMAD GHODDOSI

BIOLOGICAL NEURON










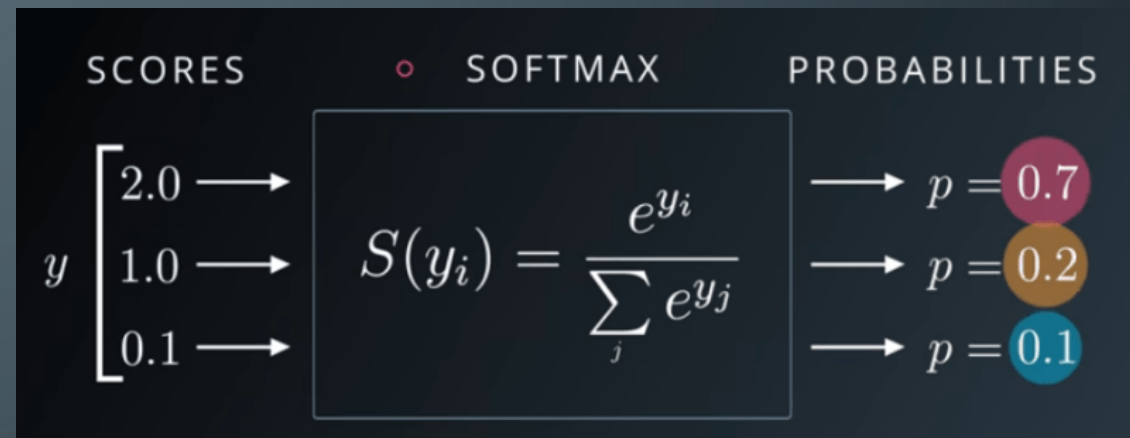
COMPUTATIONAL NEURON



ACTIVATION FUNCTIONS



Name	Plot	Equation
Identity		$f(x) = x$
Binary step		$f(x) = \begin{cases} 0 & \text{for } x < 0 \\ 1 & \text{for } x \geq 0 \end{cases}$
Logistic (a.k.a Soft step)		$f(x) = \frac{1}{1 + e^{-x}}$
TanH		$f(x) = \tanh(x) = \frac{2}{1 + e^{-2x}} - 1$
ArcTan		$f(x) = \tan^{-1}(x)$



PERCEPTRON



- Old algorithm (1958)
- Perceptron is a basic algorithm for neural networks
- Much like logistic regression
- linear problems
- Hebbian learning rule

PERCEPTRON OUTPUT



- Activation function = step function

- $$f(z) = \begin{cases} 1 & z > 0 \\ -1 & z \leq 0 \end{cases}$$



PERCEPTRON LEARNING RULE



- Learning rate: $\alpha=1$
- Perceptron input: $x^{(i)}$ and label: $y^{(i)}$
- Perceptron output: $h^{(i)} = f(z^{(i)})$
- If $y^{(i)} = h^{(i)}$ then do nothing
- Else
 - If $y^{(i)} = -1$ and $h^{(i)} = 1$ then $w(t+1) = w(t) - x^{(i)}$
 - If $y^{(i)} = 1$ and $h^{(i)} = -1$ then $w(t+1) = w(t) + x^{(i)}$




PERCEPTRON LEARNING RULE

- Learning rate: $\alpha=1$
- Perceptron input: $x^{(i)}$ and label: $y^{(i)}$
- Perceptron output: $h^{(i)} = f(z^{(i)})$
- If $y^{(i)} = h^{(i)}$ then do nothing
- Else
 - If $y^{(i)} = -1$ and $h^{(i)} = 1$ then $w(t+1) = w(t) - x^{(i)}$
 - If $y^{(i)} = 1$ and $h^{(i)} = -1$ then $w(t+1) = w(t) + x^{(i)}$


$$w(t+1) = w(t) + y^{(i)}x^{(i)}$$





PROBLEMS WITH PERCEPTRON





PROBLEMS WITH PERCEPTRON




PROBLEMS WITH PERCEPTRON



PROBLEMS WITH PERCEPTRON



MLP






TensorFlow

Playground



www.idueba.com



MLP PROS AND CONS



- Pros
 - Flexible
 - Both regression and classification
 - Good for nonlinear data with large number of inputs
- Cons
 - black box
 - computationally very expensive and time consuming to train
 - depend a lot on training data
 - overfitting

MLP ARCHITECTURE



- Single hidden layer is enough.
- If we have enough hidden units, we can solve every problem.
- Hidden units cant use linear (identity) activation function.

SHALLOW VS DEEP NEURAL NETWORKS



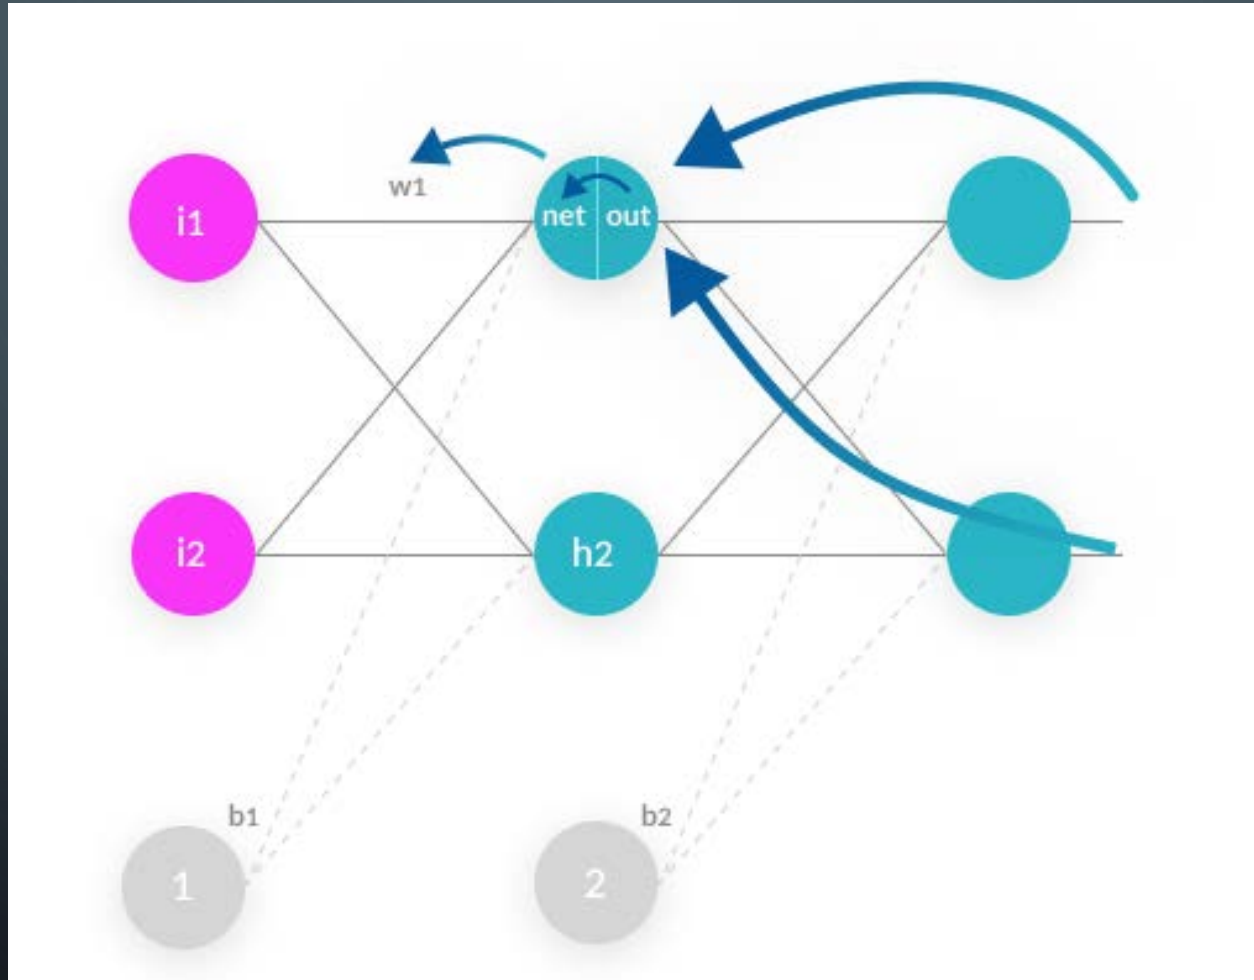
- Shallow
 - Only one hidden layer
 - Simple neurons
- Deep
 - More than one hidden layer
 - Various types of neurons
 - Convolutional
 - Recurrent
 - ...

SHALLOW VS DEEP NEURAL NETWORKS



- deep NN with the right architectures achieve better results than shallow ones
- the deep models are able to extract/build better features than shallow models

BACKPROPAGATION



OPTIMIZERS

- Gradient descent
- SGD
- mini-batch GD
- Momentum
- AdaGrad
- AdaDelta
- RMSprop
- Adam



STOCHASTIC GRADIENT DESCENT (SGD)




- Stochastic moves
- avoiding local minimum
- avoiding saddle points
- avoiding plateau
- Computational complexity

MINI-BATCH GRADIENT DESCENT



- Mini-batch
- Not too stochastic
- Fast
- Scalable
- Batch size
- epoch



CATEGORICAL CROSS-ENTROPY



- Binary cross-entropy:

- Sigmoid

$$J = \frac{1}{N} \sum_{i=1}^N -y^{(i)} * \log(h^{(i)}) - (1 - y^{(i)}) * \log(1 - h^{(i)})$$

- Categorical cross-entropy:

- Softmax

$$J = \frac{1}{N} \sum_{i=1}^N -y^{(i)} * \log(h^{(i)})$$